

Matrix Expressions and BLAS/LAPACK

Matthew Rocklin

June 27th, 2013

Need For High Level Compilers

Math is Important

// A: Naive

```
int fact(int n){  
    if (n == 0)  
        return 1;  
    return n*fact(n-1);  
}
```

// B: Programmer

```
int fact(int n){  
    int prod = n;  
    while(n--)  
        prod *= n;  
    return prod;  
}
```

// C: Mathematician

```
int fact(int n){  
    //  $n! = \text{Gamma}(n+1)$   
    return lround(exp(lgamma(n+1)))  
}
```


Evan Miller <http://www.evanmiller.org/mathematical-hacker.html>

```
x = matrix(ones(10000, 1))
```

```
x*x.T*x           Elapsed time is      ?      seconds.
```

```
(x*x.T)*x         Elapsed time is 0.337711 seconds.
```

```
x*(x.T*x)         Elapsed time is 0.000956 seconds.
```



If **A** is symmetric positive-definite and **B** is orthogonal:

Question: is $\mathbf{B} \cdot \mathbf{A} \cdot \mathbf{B}^\top$ symmetric and positive-definite?

Answer: Yes.

Question: Could a computer have told us this?

Answer: Probably.

<http://scicomp.stackexchange.com/questions/74/symbolic-software-packages-for-matrix-expressions/>

$$\beta = (X^T X)^{-1} X^T y$$

Python/NumPy

```
beta = (X.T*X).I * X.T*y
```

MatLab

```
beta = inv(X'*X) * X'*y
```

$$\beta = (X^T X)^{-1} X^T y$$

Python/NumPy

```
beta = solve(X.T*X, X.T*y)
```

MatLab

```
beta = X'*X \ X'*y
```

$$\beta = (X^T X)^{-1} X^T y$$

Python/NumPy

```
beta = solve(X.T*X, X.T*y, sym_pos=True)
```

MatLab

```
beta = (X'*X) \ (X'*y)
```


Numeric libraries for dense linear algebra

- ▶ DGEMM - **D**ouble precision **GE**neral **M**atrix **M**ultiply – $\alpha AB + \beta C$

Numeric libraries for dense linear algebra

- ▶ **DGEMM** - **D**ouble precision **GE**neral **M**atrix **M**ultiply – $\alpha AB + \beta C$
 - ▶ SUBROUTINE DGEMM(TRANSA,TRANSB,M,N,K,ALPHA,A,LDA,B,LDB,BETA,C,LDC)

Numeric libraries for dense linear algebra

- ▶ **DGEMM** - **D**ouble precision **GE**neral **M**atrix **M**ultiply – $\alpha AB + \beta C$
 - ▶ SUBROUTINE DGEMM(TRANSA,TRANSB,M,N,K,ALPHA,A,LDA,B,LDB,BETA,C,LDC)
- ▶ **DSYMM** - **D**ouble precision **SY**mmetric **M**atrix **M**ultiply – $\alpha AB + \beta C$

Numeric libraries for dense linear algebra

- ▶ **DGEMM** - **D**ouble precision **GE**neral **M**atrix **M**ultiply – $\alpha AB + \beta C$
 - ▶ SUBROUTINE DGEMM(TRANSA,TRANSB,M,N,K,ALPHA,A,LDA,B,LDB,BETA,C,LDC)
- ▶ **DSYMM** - **D**ouble precision **SY**mmetric **M**atrix **M**ultiply – $\alpha AB + \beta C$
 - ▶ SUBROUTINE DSYMM(SIDE,UPLO,M,N,ALPHA,A,LDA,B,LDB,BETA,C,LDC)

Numeric libraries for dense linear algebra

- ▶ **DGEMM** - **D**ouble precision **GE**neral **M**atrix **M**ultiply – $\alpha AB + \beta C$
 - ▶ SUBROUTINE DGEMM(TRANSA,TRANSB,M,N,K,ALPHA,A,LDA,B,LDB,BETA,C,LDC)
- ▶ **DSYMM** - **D**ouble precision **SY**mmetric **M**atrix **M**ultiply – $\alpha AB + \beta C$
 - ▶ SUBROUTINE DSYMM(SIDE,UPLO,M,N,ALPHA,A,LDA,B,LDB,BETA,C,LDC)
- ▶ ...

Numeric libraries for dense linear algebra


- ▶ **DGEMM** - **D**ouble precision **GE**neral **M**atrix **M**ultiply – $\alpha AB + \beta C$
 - ▶ SUBROUTINE DGEMM(TRANSA,TRANSB,M,N,K,ALPHA,A,LDA,B,LDB,BETA,C,LDC)
- ▶ **DSYMM** - **D**ouble precision **SY**mmetric **M**atrix **M**ultiply – $\alpha AB + \beta C$
 - ▶ SUBROUTINE DSYMM(SIDE,UPLO,M,N,ALPHA,A,LDA,B,LDB,BETA,C,LDC)
- ▶ ...
- ▶ **DPOSV** - **D**ouble symmetric **PO**sitive definite matrix **SolVe** – $A^{-1}y$

Numeric libraries for dense linear algebra

- ▶ **DGEMM** - **D**ouble precision **GE**neral **M**atrix **M**ultiply – $\alpha AB + \beta C$
 - ▶ SUBROUTINE DGEMM(TRANSA,TRANSB,M,N,K,ALPHA,A,LDA,B,LDB,BETA,C,LDC)
- ▶ **DSYMM** - **D**ouble precision **SY**mmetric **M**atrix **M**ultiply – $\alpha AB + \beta C$
 - ▶ SUBROUTINE DSYMM(SIDE,UPLO,M,N,ALPHA,A,LDA,B,LDB,BETA,C,LDC)
- ▶ ...
- ▶ **DPOSV** - **D**ouble symmetric **PO**sitive definite matrix **SolVe** – $A^{-1}y$
 - ▶ SUBROUTINE DPOSV(UPLO, N, NRHS, A, LDA, B, LDB, INFO)

Naive : $(X.T * X).I * X.T * y$

Expert : `solve(X.T*X, X.T*y, sym_pos=True)`



```
subroutine least_squares(X, y)
```

```
...
```

```
endsubroutine
```


Matrix Expressions and SymPy


SymPy Expressions

Operators (Add, log, exp, sin, integral, derivative, ...) are Python classes

Terms (3, x, log(3*x), integral(x**2), ...) are Python objects

```
>>> x = Symbol('x')
>>> expr = log(3*exp(x + 2))

>>> simplify(expr)
x + 2 + log(3)
```




SymPy Expressions

Operators (Add, log, exp, sin, integral, derivative, ...) are Python classes

Terms (3, x, log(3*x), integral(x**2), ...) are Python objects

```
>>> x = Symbol('x')
>>> expr = log(3*exp(x + 2))

>>> simplify(expr)
x + 2 + log(3)
```



Inference

```
>>> x = Symbol('x')
>>> y = Symbol('y')

>>> facts = Q.real(x) & Q.positive(y)
>>> query = Q.positive(x**2 + y)


>>> ask(query, facts)
True
```

Matrix Expressions

```
>>> X = MatrixSymbol('X', n, m)
```

```
>>> y = MatrixSymbol('y', n, 1)
```

```
>>> beta = (X.T*X).I * X.T*y
```



Matrix Inference

If **A** is symmetric positive-definite and **B** is orthogonal:

Question: is $\mathbf{B} \cdot \mathbf{A} \cdot \mathbf{B}^\top$ symmetric and positive-definite?

Answer: Yes.

Question: Could a computer have told us this?

Answer: Probably.

sympy.matrices.expressions

```
>>> A = MatrixSymbol('A', n, n)
>>> B = MatrixSymbol('B', n, n)

>>> facts = Q.symmetric(A) & Q.positive_definite(A) & Q.orthogonal(B)
>>> query = Q.symmetric(B*A*B.T) & Q.positive_definite(B*A*B.T)


>>> ask(query, facts)
True
```

Computations

Numeric libraries for dense linear algebra

- ▶ **DGEMM** - **D**ouble precision **GE**neral **M**atrix **M**ultiply – $\alpha AB + \beta C$
 - ▶ SUBROUTINE DGEMM(TRANSA,TRANSB,M,N,K,ALPHA,A,LDA,B,LDB,BETA,C,LDC)
- ▶ **DSYMM** - **D**ouble precision **SY**mmetric **M**atrix **M**ultiply – $\alpha AB + \beta C$
 - ▶ SUBROUTINE DSYMM(SIDE,UPLO,M,N,ALPHA,A,LDA,B,LDB,BETA,C,LDC)
- ▶ ...
- ▶ **DPOSV** - **D**ouble symmetric **PO**sitive definite matrix **SolVe** – $A^{-1}y$
 - ▶ SUBROUTINE DPOSV(UPLO, N, NRHS, A, LDA, B, LDB, INFO)

<http://github.com/mrocklin/computations>



```
comp = (GEMM(1, X.T, X, 0, 0)
        + GEMM(1, X.T, y, 0, 0)
        + POSV(X.T*X, X.T*y))
```

```
class GEMM(Computation):
```

```
    """ General Matrix Multiply """
```

```
    inputs      = [alpha, A, B, beta, C]
```

```
    outputs     = [alpha*A*B + beta*C]
```

```
    inplace     = {0: 4}
```

```
    fortran     = ....
```

```
class POSV(Computation):
```

```
    """ Symmetric Positive Definite Matrix Solve """
```

```
    inputs      = [A, y]
```


```
    outputs     = [UofCholesky(A), A.I*y]
```

```
    inplace     = {0: 0, 1: 1}
```

```
    condition   = Q.symmetric(A) & Q.positive_definite(A)
```

```
    fortran     = ....
```

<http://github.com/mrocklin/computations>



```
comp = (GEMM(1, X.T, X, 0, 0)
        + GEMM(1, X.T, y, 0, 0)
        + POSV(X.T*X, X.T*y))
```

```
class GEMM(Computation):
```

```
    """ Genreral Matrix Multiply """
```

```
    inputs      = [alpha, A, B, beta, C]
```

```
    outputs     = [alpha*A*B + beta*C]
```

```
    inplace     = {0: 4}
```

```
    fortran     = ....
```

```
class POSV(Computation):
```

```
    """ Symmetric Positive Definite Matrix Solve """
```

```
    inputs      = [A, y]
```

```
    outputs     = [UofCholesky(A), A.I*y]
```

```
    inplace     = {0: 0, 1: 1}
```

```
    condition   = Q.symmetric(A) & Q.positive_definite(A)
```

```
    fortran     = ....
```

```

subroutine f(X, y, var_7, m, n)
implicit none

integer, intent(in) :: m
integer, intent(in) :: n
real*8, intent(in) :: y(n)           ! y
real*8, intent(in) :: X(n, m)        ! X
real*8, intent(out) :: var_7(m)       ! 0 -> X'*y -> (X'*X)^-1*X'*y
real*8 :: var_8(m, m)                ! 0 -> X'*X
integer :: INFO                       ! INFO

call dgemm('N', 'N', m, 1, n, 1.0, X, n, y, n, 0.0, var_7, m)
call dgemm('N', 'N', m, m, n, 1.0, X, n, X, n, 0.0, var_8, m)
call dposv('U', m, 1, var_8, m, var_7, m, INFO)

RETURN
END

```

Logic Programming

TODO

- ▶ <http://github.com/mrocklin/term>
An interface for terms
Composable with legacy code via Monkey patching
Supports pattern matching via Unification
- ▶ <http://github.com/logpy/logpy>
Implements miniKanren, a logic programming language
- ▶ <http://github.com/logpy/strategies>
Partially implements Stratego, a control flow programming language

Automation

Have

```
(X.T*X).I*X.T*y  
Q.fullrank(X)
```

Want

```
comp = (GEMM(1, X.T, X, 0, 0)  
        + GEMM(1, X.T, y, 0, 0)  
        + POSV(X.T*X, X.T*y))
```


Accomplish Using Pattern Matching

```
# Source Expression, Target Computation, Condition  
(alpha*A*B + beta*C, SYMM(alpha, A, B, beta, C), Q.symmetric(A) | Q.symmetric(B  
(alpha*A*B + beta*C, GEMM(alpha, A, B, beta, C), True),  
(A.I*B, POSV(A, B), Q.symmetric(A) & Q.positive_definite(A)  
(A.I*B, GESV(A, B), True),  
(alpha*A + B, AXPY(alpha, A, B), True),
```


Pattern Matching done with LogPy, a composable Logic Programming library

$$(X^*X)^{-1}X^*y$$


#	Source Expression,	Target Computation,	Condition
	(alpha*A*B + beta*C,	SYMM(alpha, A, B, beta, C),	Q.symmetric(A) Q.symmetric(B
	(alpha*A*B + beta*C,	GEMM(alpha, A, B, beta, C),	True),
	(A.I*B,	POSV(A, B),	Q.symmetric(A) & Q.positive_definite(A)
	(A.I*B,	GESV(A, B),	True),
	(alpha*A + B,	AXPY(alpha, A, B),	True),



#	Source Expression,	Target Computation,	Condition
(alpha*A*B + beta*C, <td>SYMM(alpha, A, B, beta, C),</td> <td>Q.symmetric(A) Q.symmetric(B</td>	SYMM(alpha, A, B, beta, C),	Q.symmetric(A) Q.symmetric(B	
(alpha*A*B + beta*C, <td>GEMM(alpha, A, B, beta, C),</td> <td>True),</td>	GEMM(alpha, A, B, beta, C),	True),	
(A.I*B, <td>POSV(A, B),</td> <td>Q.symmetric(A) & Q.positive_definite(A)</td>	POSV(A, B),	Q.symmetric(A) & Q.positive_definite(A)	
(A.I*B, <td>GESV(A, B),</td> <td>True),</td>	GESV(A, B),	True),	
(alpha*A + B, <td>AXPY(alpha, A, B),</td> <td>True),</td>	AXPY(alpha, A, B),	True),	



#	Source Expression,	Target Computation,	Condition
	$(\alpha A * B + \beta C,$	$\text{SYMM}(\alpha, A, B, \beta, C),$	$Q.\text{symmetric}(A) \mid Q.\text{symmetric}(B$
	$(\alpha A * B + \beta C,$	$\text{GEMM}(\alpha, A, B, \beta, C),$	True),
	$(A.I * B,$	$\text{POSV}(A, B),$	$Q.\text{symmetric}(A) \ \& \ Q.\text{positive_definite}(A)$
	$(A.I * B,$	$\text{GESV}(A, B),$	True),
	$(\alpha A * A + B,$	$\text{AXPY}(\alpha, A, B),$	True),



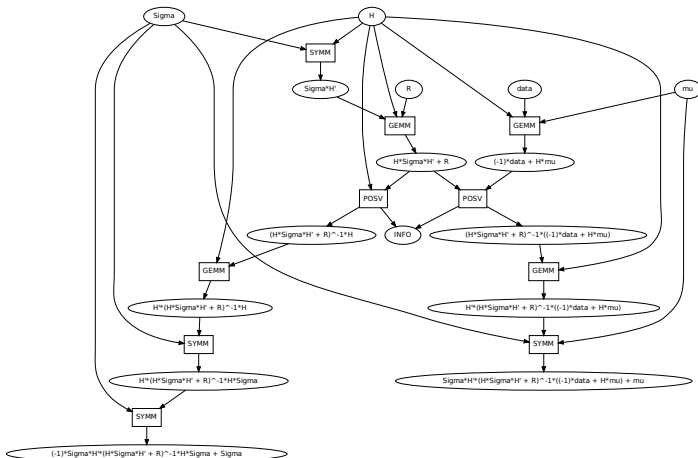
#	Source Expression,	Target Computation,	Condition
	$(\alpha A + \beta C)$	$SYMM(\alpha, A, B, \beta, C)$	$Q.symmetric(A) \mid Q.symmetric(B)$
	$(\alpha A + \beta C)$	$GEMM(\alpha, A, B, \beta, C)$	True ,
	$(A.I * B)$	$POSV(A, B)$	$Q.symmetric(A) \ \& \ Q.positive_definite(A)$
	$(A.I * B)$	$GESV(A, B)$	True ,
	$(\alpha A + B)$	$AXPY(\alpha, A, B)$	True ,

Kalman Filter

```
newmu      = mu + Sigma*H.T * (R + H*Sigma*H.T).I * (H*mu - data)
newSigma   = Sigma - Sigma*H.T * (R + H*Sigma*H.T).I * H * Sigma
```

```
assumptions = [Q.positive_definite(Sigma), Q.symmetric(Sigma),
               Q.positive_definite(R), Q.symmetric(R), Q.fullrank(H)]
```

```
f = fortran_function([mu, Sigma, H, R, data], [newmu, newSigma], *assumptions)
```



Kalman Filter

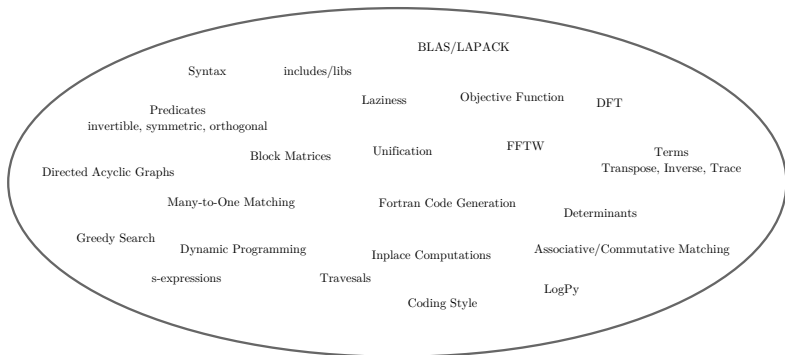
```
subroutine f(mu, Sigma, H, R, data, mu_2, Sigma_2)
  implicit none
```

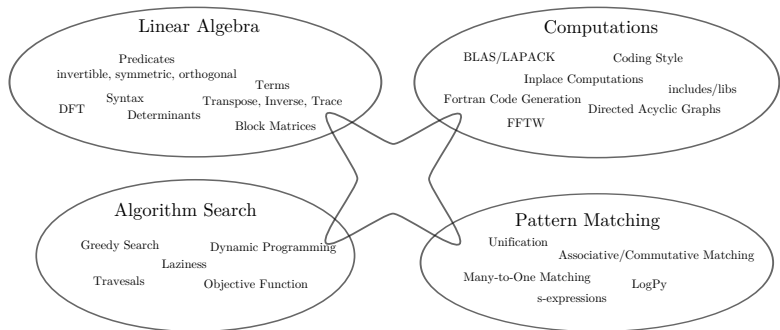
```
  real(kind=8), intent(in) :: mu(:)
  real(kind=8), intent(in) :: Sigma(:, :)
  real(kind=8), intent(in) :: H(:, :)
  real(kind=8) :: R(:, :)
  real(kind=8) :: data(:)
  real(kind=8), intent(out) :: mu_2(:)
  real(kind=8), intent(out) :: Sigma_2(:, :)
```

```
! Variable Declarations !    -- Cut for space --
```

```
  call dcopy(n**2, Sigma, 1, Sigma_2, 1)
  call dgemm('N', 'N', k, 1, n, 1.0d+0, H, k, mu, n, -1.0d+0, data, k)
  call dcopy(n, mu, 1, mu_2, 1)
  call dcopy(k*n, H, 1, H_2, 1)
  call dsymm('R', 'U', k, n, 1.0d+0, Sigma, n, H, k, 0.0d+0, var_17, k)
  call dgemm('N', 'T', k, k, n, 1.0d+0, H, k, var_17, k, 1.0d+0, R, k)
  call dposv('U', k, n, R, k, H_2, k, INFO)
  call dpotrs('U', k, 1, R, k, data, k, INFO)
  call dgemm('T', 'N', n, 1, k, 1.0d+0, H, k, data, k, 0.0d+0, var_12, n)
  call dsymm('L', 'U', n, 1, 1.0d+0, Sigma, n, var_12, n, 1.0d+0, mu_2, n)
  call dgemm('T', 'N', n, n, k, 1.0d+0, H, k, H_2, k, 0.0d+0, var_19, n)
  call dsymm('R', 'U', n, n, 1.0d+0, Sigma, n, var_19, n, 0.0d+0, var_18, n)
  call dsymm('L', 'U', n, n, -1.0d+0, Sigma_2, n, var_18, n, 1.0d+0, Sigma_2, n)
  return
```

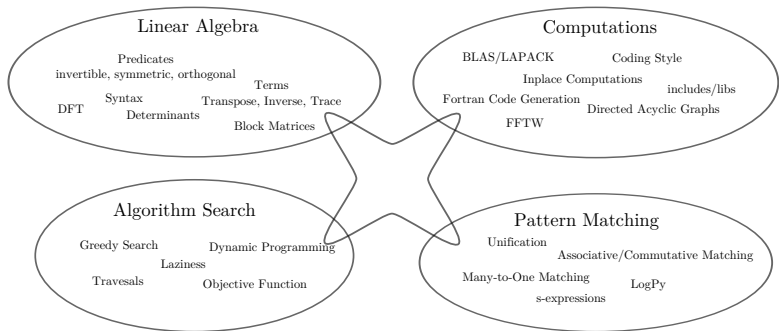
Software Design





`git@github.com:sympy/sympy.git`

`git@github.com:mrocklin/computations.git`




`git@github.com/logpy/strategies.git`

`git@github.com/logpy/logpy.git`


```
X = MatrixSymbol('X', n, m)
y = MatrixSymbol('y', n, 1)
```


```
inputs  = [X, y]
outputs = [(X.T*X).I*X.T*y]
facts   = fullrank(X)
```



```
X = MatrixSymbol('X', n, m)
y = MatrixSymbol('y', n, 1)
```

```
inputs  = [X, y]
outputs = [(X.T*X).I*X.T*y]
facts   = fullrank(X)
```





```

class SYRK(Computation):
    """ Symmetric Rank-K Update 'alpha X' X + beta Y' """
    inputs = (alpha, A, beta, D)
    outputs = (alpha * A * A.T + beta * D,)
    inplace = {0: 3}
    fortran_template = ("call %(fn)s('%(UPLO)s', '%(TRANS)s', %(N)s, %(K)s, "
                                "%(alpha)s, %(A)s, %(LDA)s, "
                                "%(beta)s, %(D)s, %(LDD)s)")
    ...





(alpha*A*A.T + beta*D, SYRK(alpha, A, beta, D), True),
(A*A.T, SYRK(1.0, A, 0.0, 0), True),

```

Separation promotes Comparison and Experimentation

`git@github.com:sympy/sympy.git`

`git@github.com:mrocklin/computations.git`




`git@github.com/logpy/strategies.git`




`git@github.com/logpy/logpy.git`

Separation promotes Comparison and Experimentation

`git@github.com:sympy/sympy.git`



`git@github.com:Theano/Theano.git`

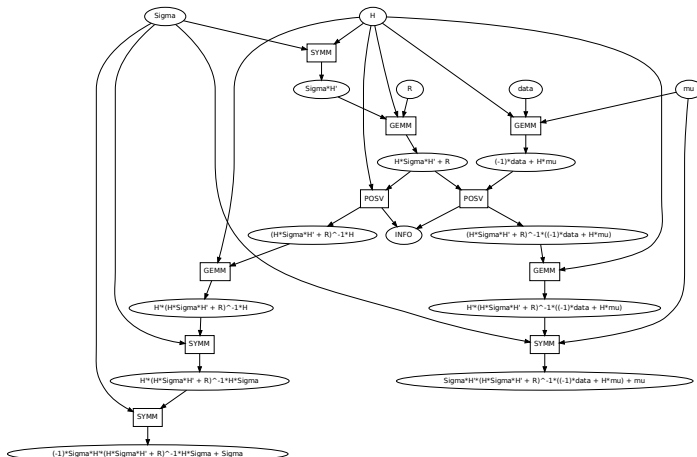


Kalman Filter - Theano v. Fortran

```
newmu      = mu + Sigma*H.T * (R + H*Sigma*H.T).I * (H*mu - data)
newSigma   = Sigma - Sigma*H.T * (R + H*Sigma*H.T).I * H * Sigma
```

```
assumptions = [Q.positive_definite(Sigma), Q.symmetric(Sigma),
               Q.positive_definite(R), Q.symmetric(R), Q.fullrank(H)]
```


```
f = fortran_function([mu, Sigma, H, R, data], [newmu, newSigma], *assumptions)
```



Kalman Filter - Theano v. Fortran

```
newmu      = mu + Sigma*H.T * (R + H*Sigma*H.T).I * (H*mu - data)
newSigma   = Sigma - Sigma*H.T * (R + H*Sigma*H.T).I * H * Sigma
```

```
f = theano_function( [mu, Sigma, H, R, data], [newmu, newSigma])
```






Final Thoughts

- ▶ Modularity is good!
 - ▶ Cater to single-field experts
 - ▶ Eases comparison and evolution
 - ▶ This project might die but the parts will survive
- ▶ Intermediate Representations are Good!
 - ▶ Fortran code doesn't depend on Python
 - ▶ Readability encourages development
 - ▶ Extensibility (lets generate CUDA)
- ▶ Read more! <http://matthewrocklin.com/blog>
- ▶ Listen more!
 - ▶ Dynamics with SymPy Mechanics, *Jason Moore*, Room 204 - 2:10pm
 - ▶ SymPy Gamma and SymPy Live: Python and Mathematics Online, *David Li* Room 203 - 3:50pm

Multiple Results

$$(X^*X)^{-1}X^*y$$



Multiple Results

$$(X^*X)^{-1}X^*y$$

