

Import without a filesystem:

scientific Python built-in with static linking
and frozen modules

Pat Marion, Kitware Inc.

Aron Ahmadia

Bradley M. Froehle, University of California, Berkeley

Presentation outline

- Intro
- Problem statement
- Summary of approaches
- This talk's approach: built-in modules
- Challenges, automation, adoption

About me

- Pat Marion
- R&D Engineer @ Kitware
- Lately, I do point cloud processing with python/c++
- Previously, supercomputing with python

Talk accompanying material

github.com/patmarion/NumpyBuiltinExample

Problem statement

Python *import* scales very poorly when parallel process independently request filesystem metadata from a shared, network filesystem

How bad is it?

“catastrophic”

-William Scullin “**Python for High Performance Computing**” PyCon 2011.

“For 32k processes on BlueGene/P...

35 minutes to load and initialize interpreter

5.5 hours to import 100 trivial C-extension modules”

-Asher Langton “**Improving Python+MPI import performance**” Jan. 2012 numpy-discussion.

Aron's talk last year

Solving the Import Problem: scalable dynamic loading

<http://pyvideo.org/video/1201/solving-the-import-problem-scalable-dynamic-load>

How import works

Brett Cannon's talk from PyCon 2013:

<http://pyvideo.org/video/1707/how-import-works>

“how the simple `from sys import version` turns out to be slightly complicated”

Import documentation

- PEP 302 --- New Import Hooks
<http://www.python.org/dev/peps/pep-0302/>
- Python 3 Language Reference: The Import System
<http://docs.python.org/3/reference/import.html>

import numpy

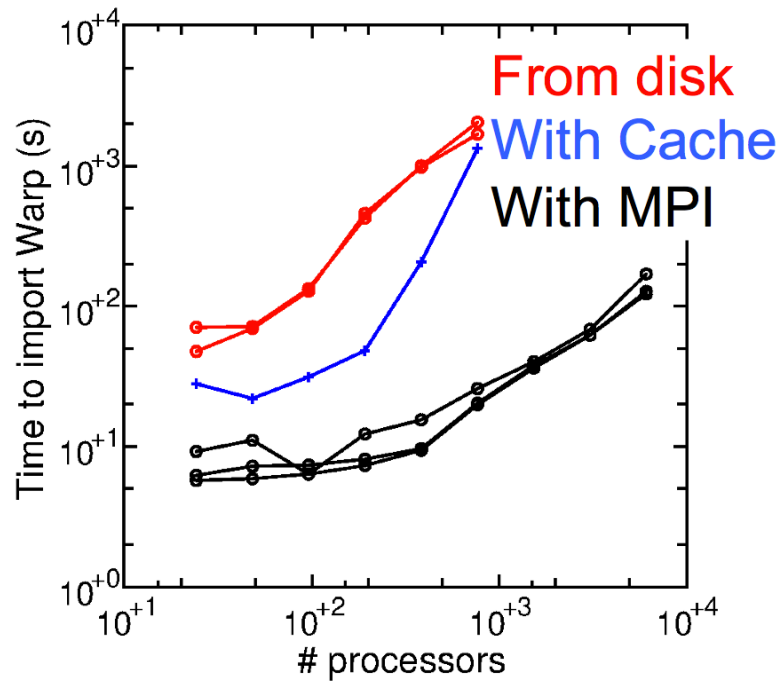
- Search each directory listed in *sys.path*
- Read python module (.py, .pyc)
- Read C extension module (.so)

```
>>> import numpy
```

```
>>> len([m for m in sys.modules.keys() if 'numpy' in m])
```

```
140
```

Timing results from Hopper



Results from Dave Grote's talk *Python in a Parallel Environment* @ NERSC User Group meeting 2013

Two ways to fix

- Change Python's import mechanism
- Change filesystem access methods

Summary of approaches

- Intercept basic I/O ops and replace with MPI
- Path caching
- Built-in

Projects

- walla, collfs, DLcache, FMcache
- mpi_import.py, cached_import.py
- Slither

Today's topic

- Built-in: import without a filesystem
- Combination of build techniques:
 - built-in modules
 - frozen modules

Experiment at Argonne National Labs

- Intrepid IBM Blue Gene/P @ ALCF
- Ran PHASTA simulation code with ParaView Python coprocessor @ 160k cores (full machine scale)
- Without import solution, @ 32k nodes our job was terminated before import completed (>1 hour)
- With built-in + frozen modules, import took <0.1 seconds at @ 160k cores.

What's a built-in module?

```
>>> import sys
```

```
>>> sys.builtin_module_names
```

```
('__builtin__', '__main__', '_ast', '_codecs', '_sre',  
'_symtable', '_warnings', '_weakref', 'errno',  
'exceptions', 'gc', 'imp', 'marshal', 'posix', 'pwd',  
'signal', 'sys', 'thread', 'xxsubtype', 'zipimport')
```


module.__file__

```
>>> import math
```

```
>>> math.__file__
```

```
‘/usr/lib/python2.7/lib-dynload/math.so’
```

Math built-in

```
>>> import math
```

```
>>> math
```

```
<module 'math' (built-in)>
```

```
>>> math.__file__
```

```
Traceback (most recent call last):
```

```
  File "<stdin>", line 1, in <module>
```

```
AttributeError: 'module' object has no attribute '__file__'
```

built-in: how to

- Place Setup.local in the Python Modules/ source code directory. Setup.local lists each module that should be built-in and the .c file for the module.
- Modules are compiled into libpython and linked into python interpreter.
- Python standard library modules built-in: OK
- Third party modules: hmm... what? I have to recompile libpython? No, but you must recompile the python interpreter and link your module library, or create your own python interpreter.

built-in: how it works

- “built-in” only works for C extension modules
- Register “math” with pointer to `initmath()` function.
- The `initmath()` function registers the module methods. The “import math
- When importing, Python’s import machinery looks at the table of built-in module names, finds “math” and calls `initmath()` directly
- ...instead of `dlopen(“math.so”); dlsym(“initmath”);`

Why is a patch required?

- Python's support for built-in modules does not work for extension modules that are sub-modules (extensions inside a module package)
- Example of an extension module that is a sub-module:

```
>>> numpy.core.multiarray.__file__
```

```
‘/usr/lib/dist-packages/numpy/core/multiarray.so’
```

Python Freeze for .py files

- You can embed .py files using the freeze tool
- Part of Python source code distribution: Tools/freeze/freeze.py
- `$ freeze.py myscript.py`
- `myscript.py` must be importable
- Collects list of all imported .py files
- Writes a C source file with the .py file content as a string (char array)

Python Freeze for .py files

```
>>> import os
```

```
>>> os
```

```
<module 'os' from '<frozen>'>
```

```
>>> os.__file__
```

```
'<frozen>'
```

Python Freeze for .py files

```
>>> import imp
```

```
>>> imp.is_frozen('os')
```

```
True
```


Python Freeze for .py files

- Just like built-in modules, you must recompile the python interpreter to add the frozen code source files (or create your own interpreter)
- Just like built-in modules, Python will search for module names in the table of frozen modules, it will get the module content as a char array instead of `fopen(myfile.py); fread(file);`

Remember: the main problem is the metadata

By using built-in and frozen modules, Python can locate modules by searching module tables in memory. No more searching directories in `sys.path`

Built-in + freeze: Pros

- It's really fast – avoids filesystem access
- Built-in and freeze is officially supported by Python, no hacks
- It can be tested on any platform, as a single process or any scale
- Avoids dynamic loading, fully self-contained interpreter
- Although the build is complicated, once you have that figured out, then everything “just works”

Built-in + freeze: Cons

- Must compile a custom interpreter
- Requires a patch to Python (shouldn't have to...)
- Must recompile if you change a .py file
- Build method to add third party modules is more complicated, sometimes very complicated.

Slither: static Python builds for HPC systems

- a command line tool for building static CPython binaries
- github.com/bfroehle/slither

NumPy Built-in Example

An example project demonstrating a simple hello world C program that contains NumPy and the Python standard library, built-in and frozen:

github.com/patmarion/NumpyBuiltinExample

numpy modules

```
>>> numpy.__file__
```

```
‘/usr/lib/dist-packages/numpy/__init__.pyc’
```

```
>>> numpy.__path__
```

```
[‘/usr/lib/dist-packages/numpy’]
```

Other issues: auxiliary data

As mentioned in PEP 302 #open-issues:

“Modules often need supporting data files to do their job, particularly in the case of complex packages or full applications. Current practice is generally to locate such files via `sys.path` (or a `package.__path__` attribute). This approach will not work, in general, for modules loaded via an import hook.”

numpy.test()

- Doesn't work because nose searches for tests in the filesystem using module.__file__ attribute

Python CMake build system

github.com/davidsansome/python-cmake-buildsystem

- Portable and adaptable

Questions?