

Using IPython Notebook with IPython Cluster for Reproducibility and Portability of Atomistic Simulations

Zachary T. Trautt*, L.H. Friedman*, C.A. Becker†

* Materials Measurement Science Division,

† Materials Science and Engineering Division,


*† National Institute of Standards and Technology

Outline

- Quick introduction to my work
- Why workflow is necessary
- How I use IPython Notebook and IPython Cluster

Example: Material Structure-Properties

- Mechanical properties (strength of metal in your car)
 - Determined by microscopic structure
 - Study processing origin of microscopic structure with atomic-scale simulation

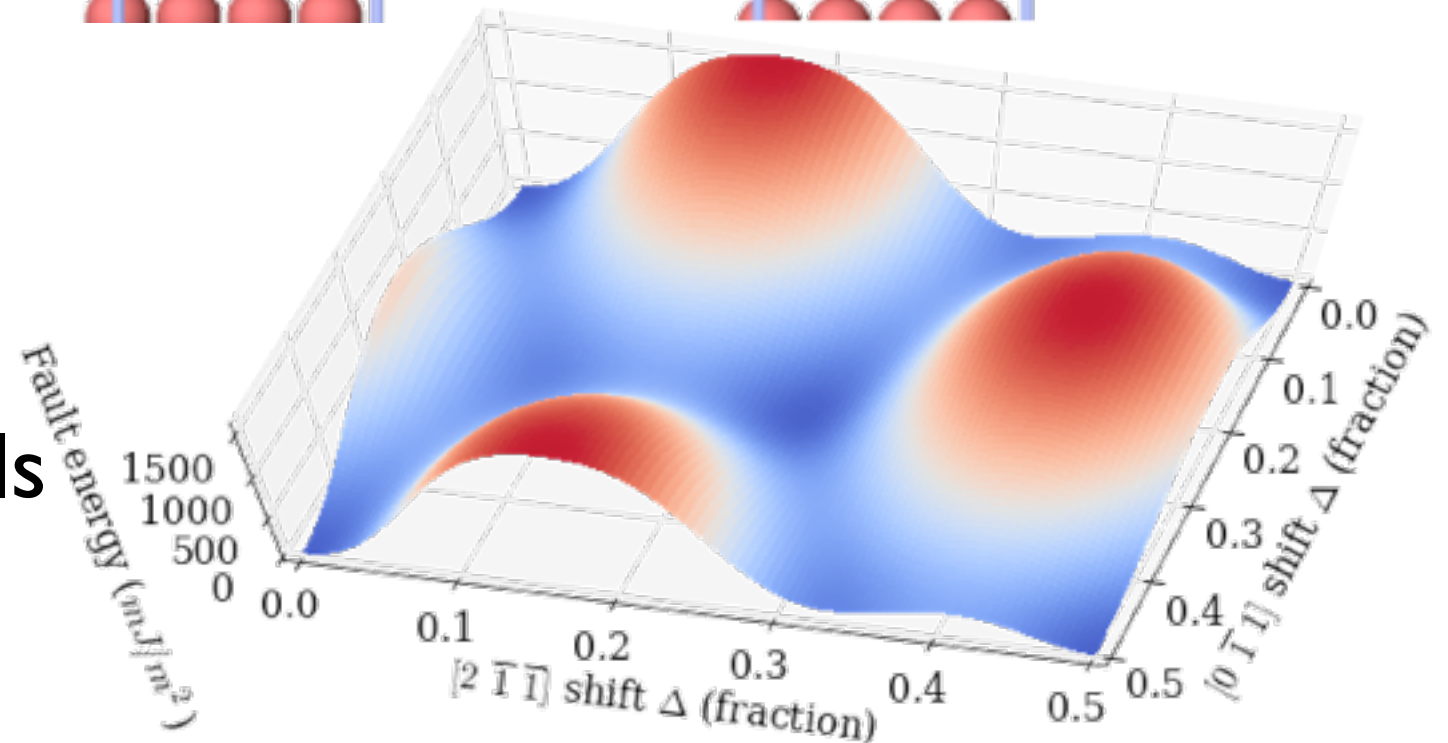



[Trautt ZT, Adland A, Karma A, Mishin Y. Acta Mater 2012;60:6528.](#)

Problem Complexity

- Many parameters: **many simulations**

- Composition
- Geometry
- Temperature
- Random seed
- Methodology
- Implementation details
- Etc.



- Large simulation domain: **parallelism mandatory**
- Solution: must address distributable parallel tasks

What I have done in the past

- A **lot** of bash scripting...

- for f in \$params; do mkdir ...
- for f in \$folders; do cd \$f; qsub preprocessing ...
- for f in \$folders; do cd \$f; qsub production ...
- for f in \$folders; do cd \$f; qsub postprocessing ...
- for f in \$folders; do gather data ...

- Reproducibility with logging tool?

- for f in \$params; do mkdir ...
- for f in \$folders; do cd \$f; qsub TheLoggerTool preprocessing ...
- for f in \$folders; do cd \$f; qsub TheLoggerTool production ...
- for f in \$folders; do cd \$f; qsub TheLoggerTool postprocessing ...
- for f in \$folders; do gather data ...

- Reproducible, but not scalable to full problem complexity.

- Reproducible tasks, but have I recorded the recipe?

- Workflow tool?

Workflow, because logging is not enough

	Workflow Tool	Log Tool
Task automation, assembly, distribution	✓ Yes	?
Reproducibility	✓ Run workflow again (same input)	✓ Rerun a task
Reuse (explore parameter space)	✓ Run workflow again (different input)	?
Benchmark on local cluster, exec. in cloud	✓ Run workflow again (more/larger/longer simulations)	?
Efficiencies gained in automation	✓ Yes	?
Easy to build upon past work	✓ Change applicable portions of workflow and run again	?

A look at traditional workflow tools, alternatives, etc.

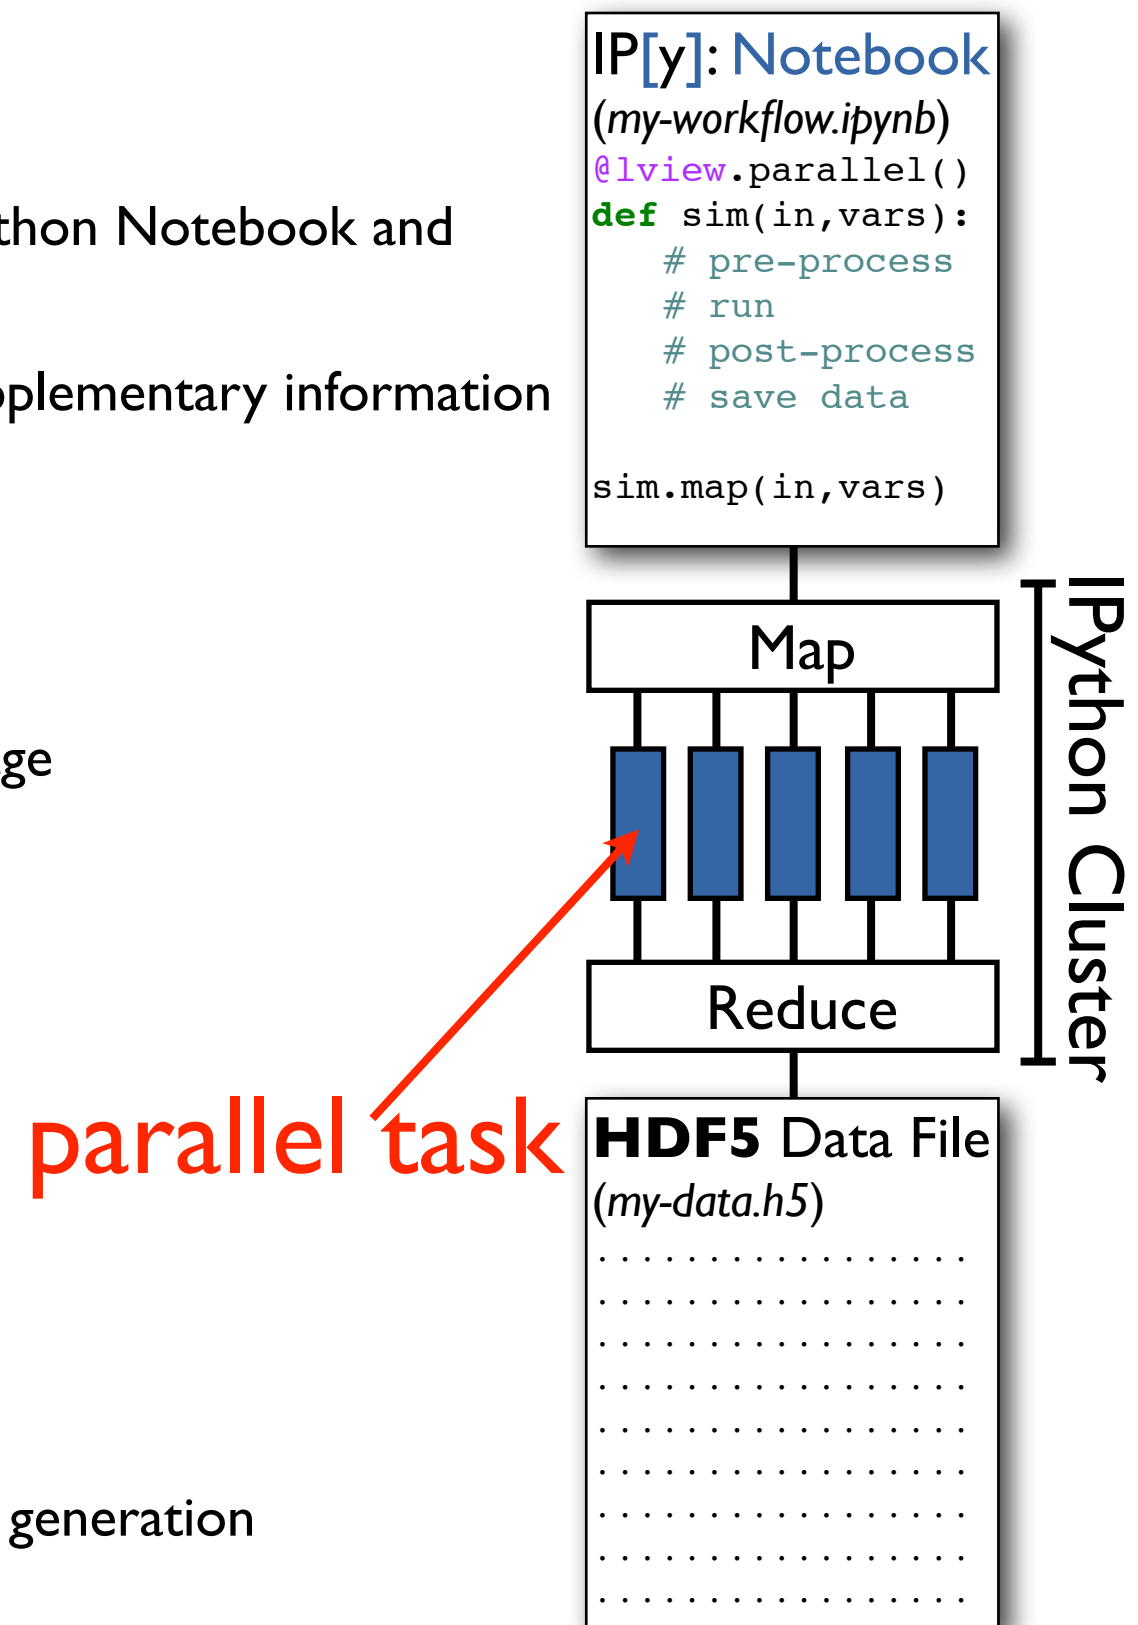
- Is it possible to distribute parallel tasks to remote resources?
 - Some: (✓) yes
 - Some: (✗) no
 - Many: (?) unknown
- Does it appear to provide more benefit (automation) than cost (time invested in installation/setup/use)?
 - Some: (✓) yes
 - Some: (✗) no
 - Many: (?) insufficient applicable examples
- Is there strong package development to internalize scientific work and avoid reinventing the wheel?
 - Some: (✓) yes
 - Most: (✗) no

A “short” list:

DAGMan, Pegasus, Triana, ICENI, Taverna, GridAnt, GrADS, GridFlow, Unicore, Gridbus workflow, Askalon, Karajan, Kepler, Grid Mathematica, Galaxy, VisTrails, Workflow4ever, WorkWays, LabArchives, iRODS, Accelrys Pipeline Pilot, RapidMiner, LONI Pipeline, GWorkflowDL, Bioclipse Scripting Language, BioExtract Server, Tesla, BioVeL, SHIWA, HTCCondor, Trident, Amazon Simple Workflow Service, Nipype, Sumatra, Wakari, Anaconda, joblib, Trac, gridfields

IPython benefits

- Easy to use!
- Plethora of packages and active development
- My entire workflow is defined and executed from IPython Notebook and results stored in a pytable
 - Both can be easily posted, shared, or added as supplementary information with a paper
- Improve reproducibility:
 - How to reproduce results
 - Download/compile molecular dynamics package
 - Install/configure IPython Cluster/Notebook
 - Run Notebook
 - Only trivial changes in notebook:
/users/remote/project/directory/name/
- End-to-end automation
 - Improve efficiency
 - Reduce random mistakes
 - Enable moving work to the “cloud” for rapid data generation



Workflow - IPython Notebook

```
I def simStep1(...):  
    import numpy  
    # do stuff  
I dview.push(dict(simStep1=simStep1));
```

← Define Tasks

```
I @lview.parallel()  
I def simulation(Temp, seed, Theta, Phi):  
    simStep1(...)  
    simStep2(...)  
    simStep3(...)
```

← Define Workflow

```
I simulation.map(TEMP, SEED, THETA, PHI)
```

← Execute

I You can do a workflow in serial

I Or distribute to cluster

TEMP=[0,100,200,300,...] # (K)

Workflow - IPython Notebook

```
def simStep1(...):  
    import numpy  
    # do stuff  
dview.push(dict(simStep1=simStep1));
```

```
@lview.parallel()  
def simulation(Temp, seed, Theta, Phi):  
    simStep1(...)  
    simStep2(...)  
    simStep3(...)
```

```
simulation.map(TEMP, SEED, THETA, PHI)
```

Direct and load-
leveled view of our
IPython Cluster

```
from IPython.parallel import Client  
rc = Client()  
lview = rc.load_balanced_view()  
lview.block = True  
dview = rc[:]  
dview.block = True
```

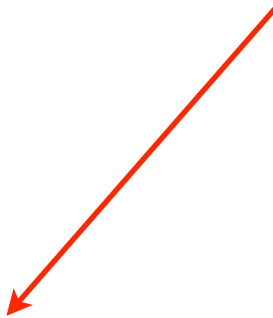
Workflow - IPython Notebook

```
def simStep1(...):  
    import numpy  
    # do stuff  
dview.push(dict(simStep1=simStep1));
```

```
@lview.parallel()  
def simulation(Temp, seed, Theta, Phi):  
    simStep1(...)  
    simStep2(...)  
    simStep3(...)
```

```
simulation.map(TEMP, SEED, THETA, PHI)
```

Make a unique
working directory



```
def mkWorkDir(basedir):  
    import os, uuid  
    uniqueID=str(uuid.uuid4())  
    workDir=basedir+'ipengine-'+uniqueID  
    os.mkdir(workDir)  
    return workDir  
dview.push(dict(mkWorkDir=mkWorkDir));
```


Workflow - IPython Notebook

```
def simStep1(...):  
    import numpy  
    # do stuff  
dview.push(dict(simStep1=simStep1));
```

```
@lview.parallel()  
def simulation(Temp, seed, Theta, Phi):  
    simStep1(...)  
    simStep2(...)  
    simStep3(...)
```

```
simulation.map(TEMP, SEED, THETA, PHI)
```

Large-scale Atomic/Molecular
Massively Parallel Simulator

<http://lammps.sandia.gov/index.html>

[S Plimpton J Comp
Phys, 1995;117:1-19.](#)

Write a LAMMPS
input script

```
def lmpscript(... timestep, ...):  
    ...  
    f.write('timestep %f\n' % timestep)  
    ...  
    f.close()  
dview.push(dict(lmpscript=lmpscript));
```

Workflow - IPython Notebook

```
def simStep1(...):  
    import numpy  
    # do stuff  
dview.push(dict(simStep1=simStep1));
```

```
@lview.parallel()  
def simulation(Temp, seed, Theta, Phi):  
    simStep1(...)  
    simStep2(...)  
    simStep3(...)
```

```
simulation.map(TEMP, SEED, THETA, PHI)
```

Large-scale Atomic/Molecular
Massively Parallel Simulator

<http://lammps.sandia.gov/index.html>

[S Plimpton J Comp
Phys, 1995;117:1-19.](#)

Run LAMMPS




```
def runlmp(workDir, lmpExe, lmpName, cpus):  
    import os, subprocess  
    os.chdir(workDir)  
    string='mpirun -np '+str(cpus)+' '+lmpExe+' < '+lmpName  
    subprocess.call([string], shell=True)  
dview.push(dict(runlmp=runlmp));
```

Parallel External Tasks - IPython Cluster

```
simulation.map(TEMP, SEED, THETA, PHI)
```


IPython Cluster for Serial Tasks:

□ : CPU core
■ : IPython Engine



More CPUs and IPengines on “cloud”

IPython Cluster for Parallel External Tasks:



Saving Data - HDF5

```
from IPython.parallel import Client
rc = Client()
myEngines=[ i for i in xrange(1,len(rc.ids))]

lview = rc.load_balanced_view(myEngines)
lview.block = True

dview = rc[1:]
dview.block = True


hview = rc[0]
hview.block = True

@dview.remote()
def sethview():
    from IPython.parallel import Client
    rc = Client()
    global hview
    hview = rc[0]
    hview.block = True
sethview();
```

```
@lview.parallel()
def simulation(...):
    ...
    hview.apply_sync(recordResult,...)

simulation.map(...)
```

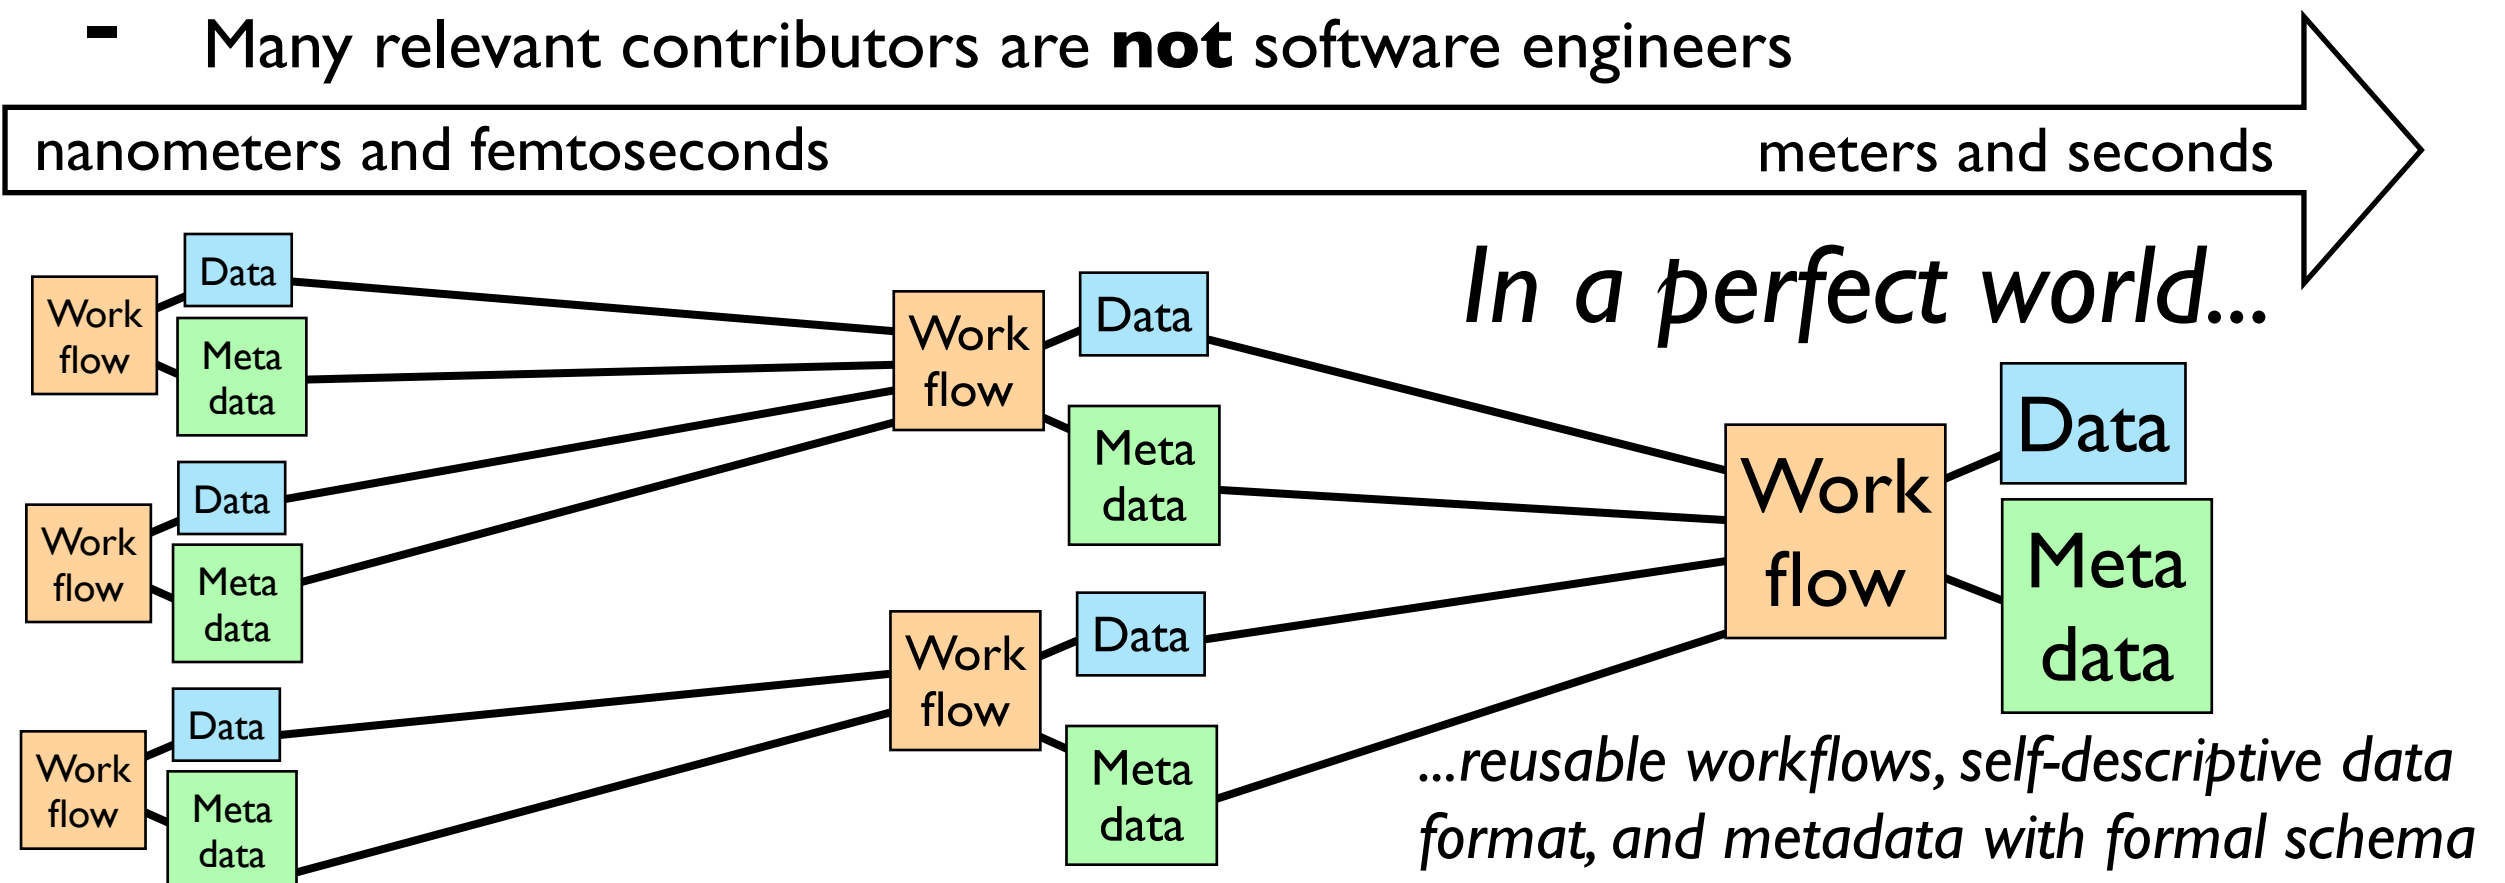
```
def recordResult(C11,...):
    ...
    simulation['C11']=C11
    simulation.append()
    ...
dview.push(dict(recordResult=recordResult));
hview.push(dict(recordResult=recordResult));
```



Example IPython notebook

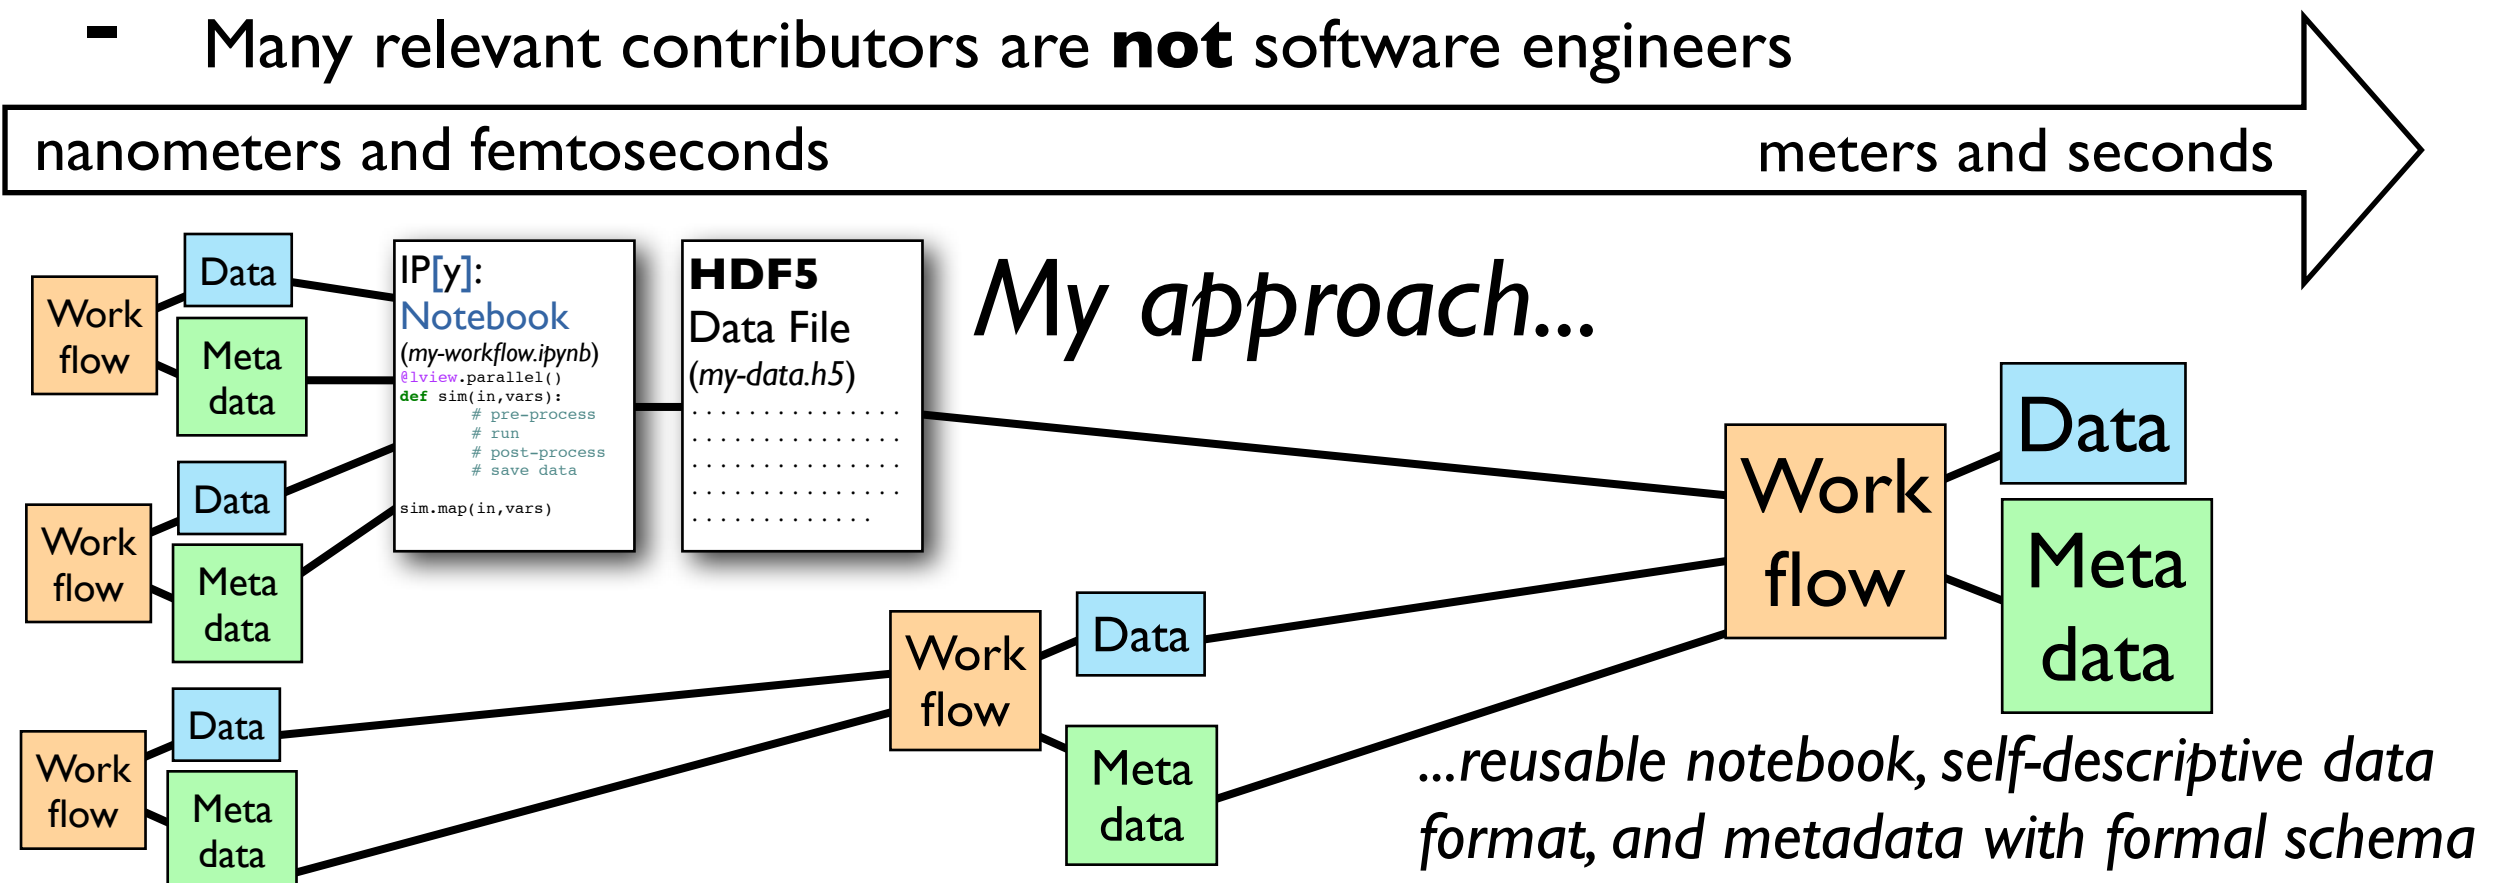
Summary

- IPython Notebook + IPython Cluster
 - Easy: use, reproducibility, reuse, etc.
 - Do more work with less effort and very little code
- Grand challenges in materials science
 - Complex problem: big, disparate data
 - Relevant behavior at different length/time scales: multi-scale modeling
 - Many relevant contributors are **not** software engineers




Summary



- IPython Notebook + IPython Cluster
 - Easy: use, reproducibility, reuse, etc.
 - Do more work with less effort and very little code
- Grand challenges in materials science
 - Complex problem: big, disparate data
 - Relevant behavior at different length/time scales: multi-scale modeling
 - Many relevant contributors are **not** software engineers



Thank you!



Atomistic Simulation



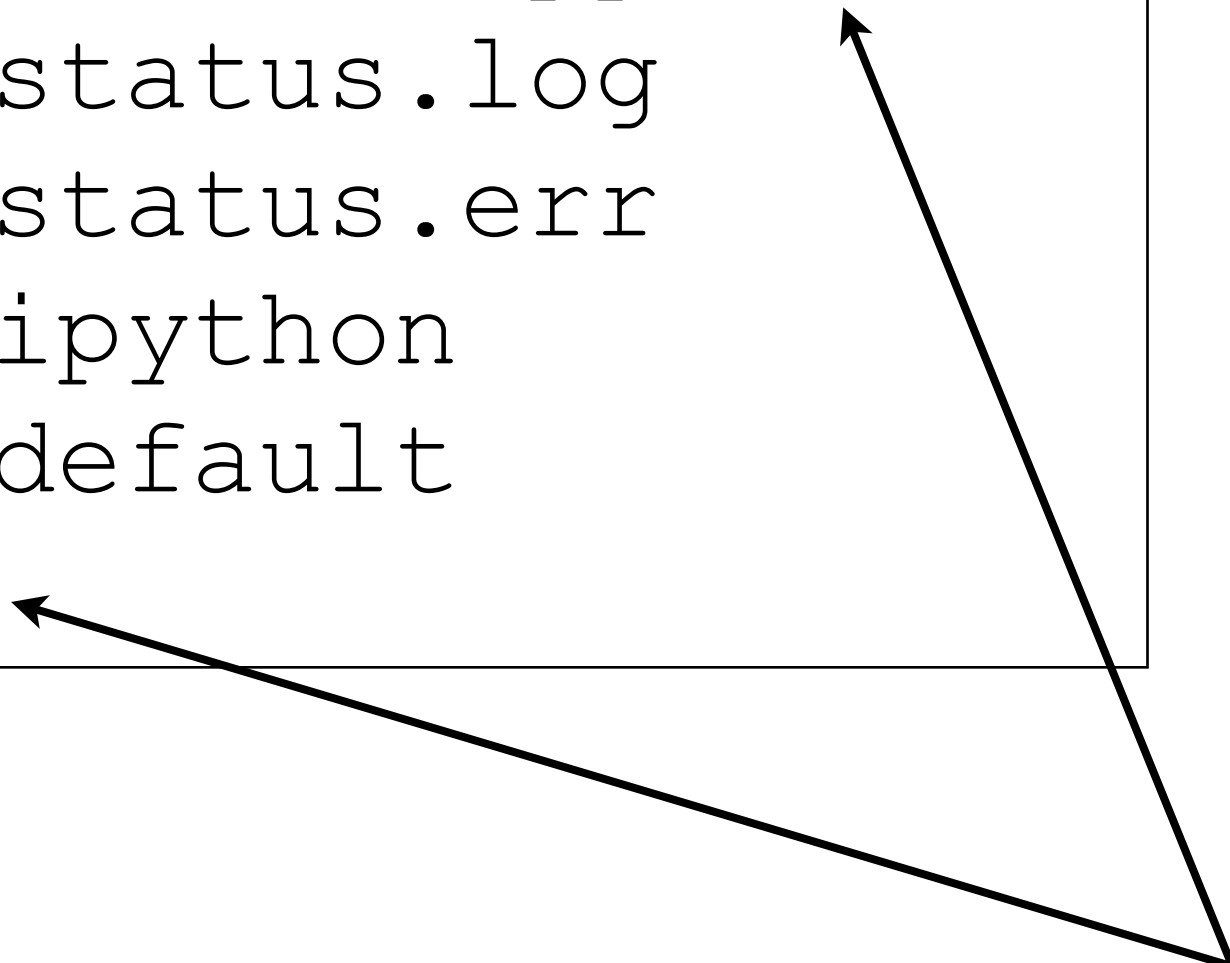
- Define an atomic configuration
- Define interaction forces and energy between atoms
- Numerically solve Newton's equation of motion

$$\vec{F} = m\vec{a}$$

IPython Engine - MPI task

Portable Batch System (PBS) Script:

```
#PBS -l nodes=1:ppn=8:amd
#PBS -o status.log
#PBS -e status.err
#PBS -N ipython
#PBS -q default
ipengine
```



8 CPUs per IPython Engine